

Student Presentation

Reproducing Egalitarian Paxos



COS 518: *Advanced Computer
Systems*


Jinzheng Tu, Yu Zeng
4/17/2019

Problem Statement / Motivation

1. Paxos:
High **latency** and low **throughput**.
2. Multi-paxos:
Leader becomes the bottleneck for **availability**.

Key Idea

- Every replica is the leader for the requests it received.
- Latency and throughput
 - Replicas do not contend
 - Can choose any fast quorum
- Availability
 - workload balanced across all replicas



Key Challenges

- Multiple leaders for multiple requests
 - Keep the dependency relationships between commands coherent across all replicas
 - Execute commands in dependency order

Key Result (Evaluation)

→ epaxos \$./compose 3 --prod up -d

Creating epaxos-server-0 ... done

Creating epaxos-server-1 ... done

Creating epaxos-server-2 ... done

PUT <server> <key> <value>

→ epaxos \$./bin/client -n 3 put -v 0 114 514

PUT: [114]=514 to 0 start

PUT: RequestMsg 4972707484880603879 sent to 0

PUT: [114]=514 to 0 committed

GET <server> <key>

→ epaxos \$./bin/client -n 3 get 1 114

<!NOT IMPLEMENTED!>

Key Result (Evaluation)

1000 random PUTs to 1000 keys on 3 servers

→ epaxos \$./bin/client -n 3 batch-put 0 1000
-N 1000 --pipeline 10000 -S 3 --random-key
--random-server >/dev/null && echo Good

Good

Key Result (Evaluation)

Emulate 10% package loss

→ epaxos \$ make pumba-loss-10 && ...

Good

Emulate 1000ms delay

→ epaxos \$ make pumba-delay-1000ms && ...

Good

Random kill

→ epaxos \$ docker kill epaxos-server-0 && ...

...Nope...

Emulate 0.1% package duplication

→ epaxos \$ make pumba-dup-0.1 && ...

...Nope...

Future Plan

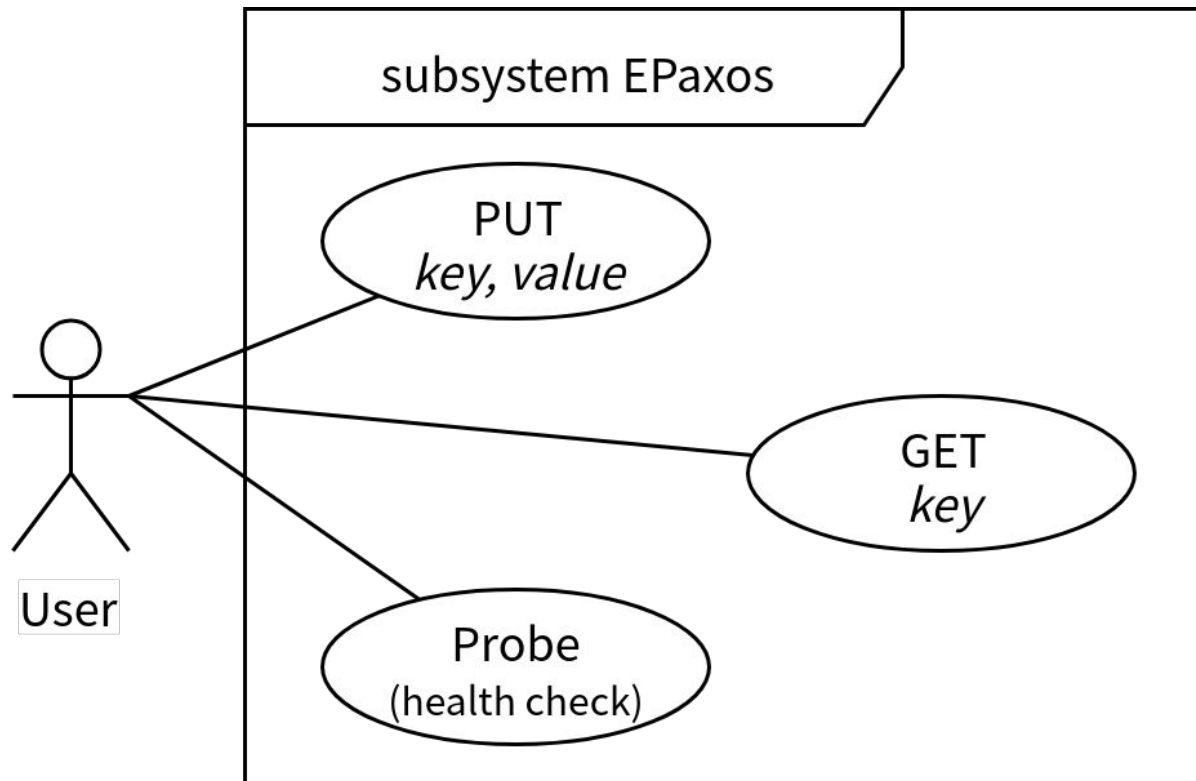
- Implement "Prepare" logic (recovery)
- Implement "GET" logic
- Debug
- and debug...
- Test correctness against (emulated) delay, duplication, loss, random kill, etc.
- Performance test by kubernetes on GKE!

Technical Details

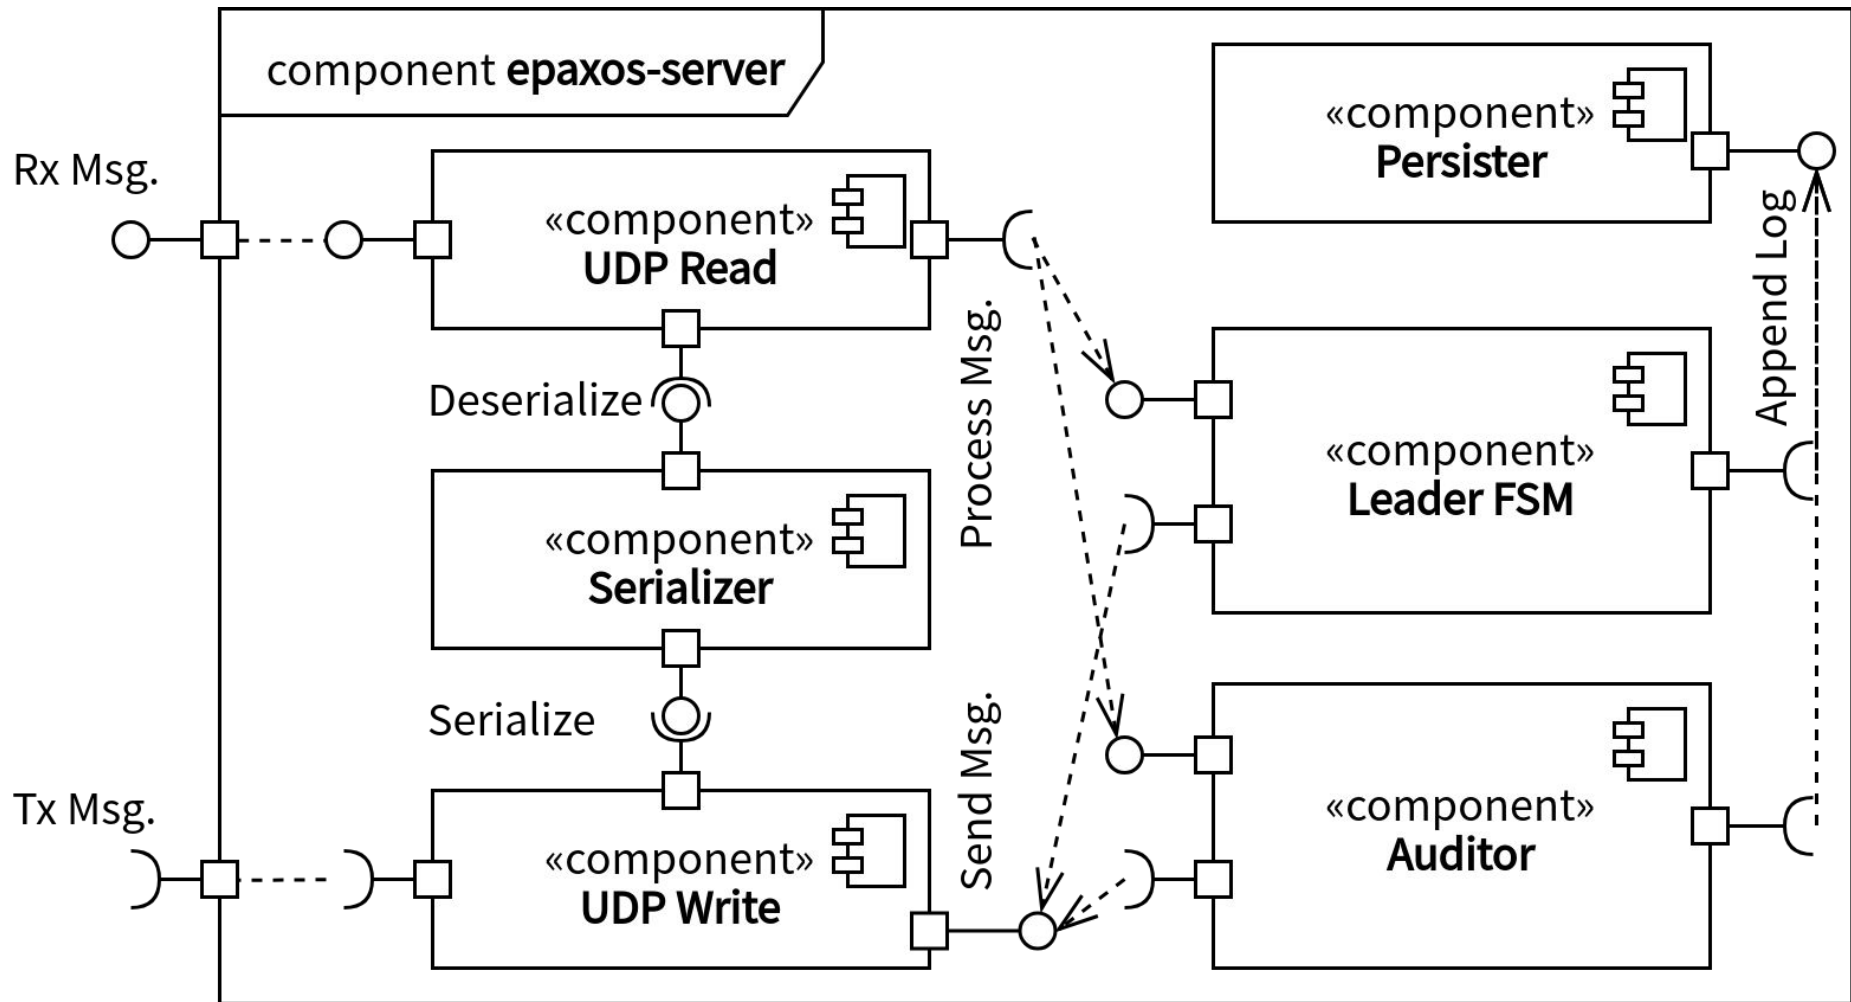
System Architecture

- "4+1" views:
 - Logical View
 - Process View
 - Development View
 - Physical View
 - ~~○ Scenario View~~

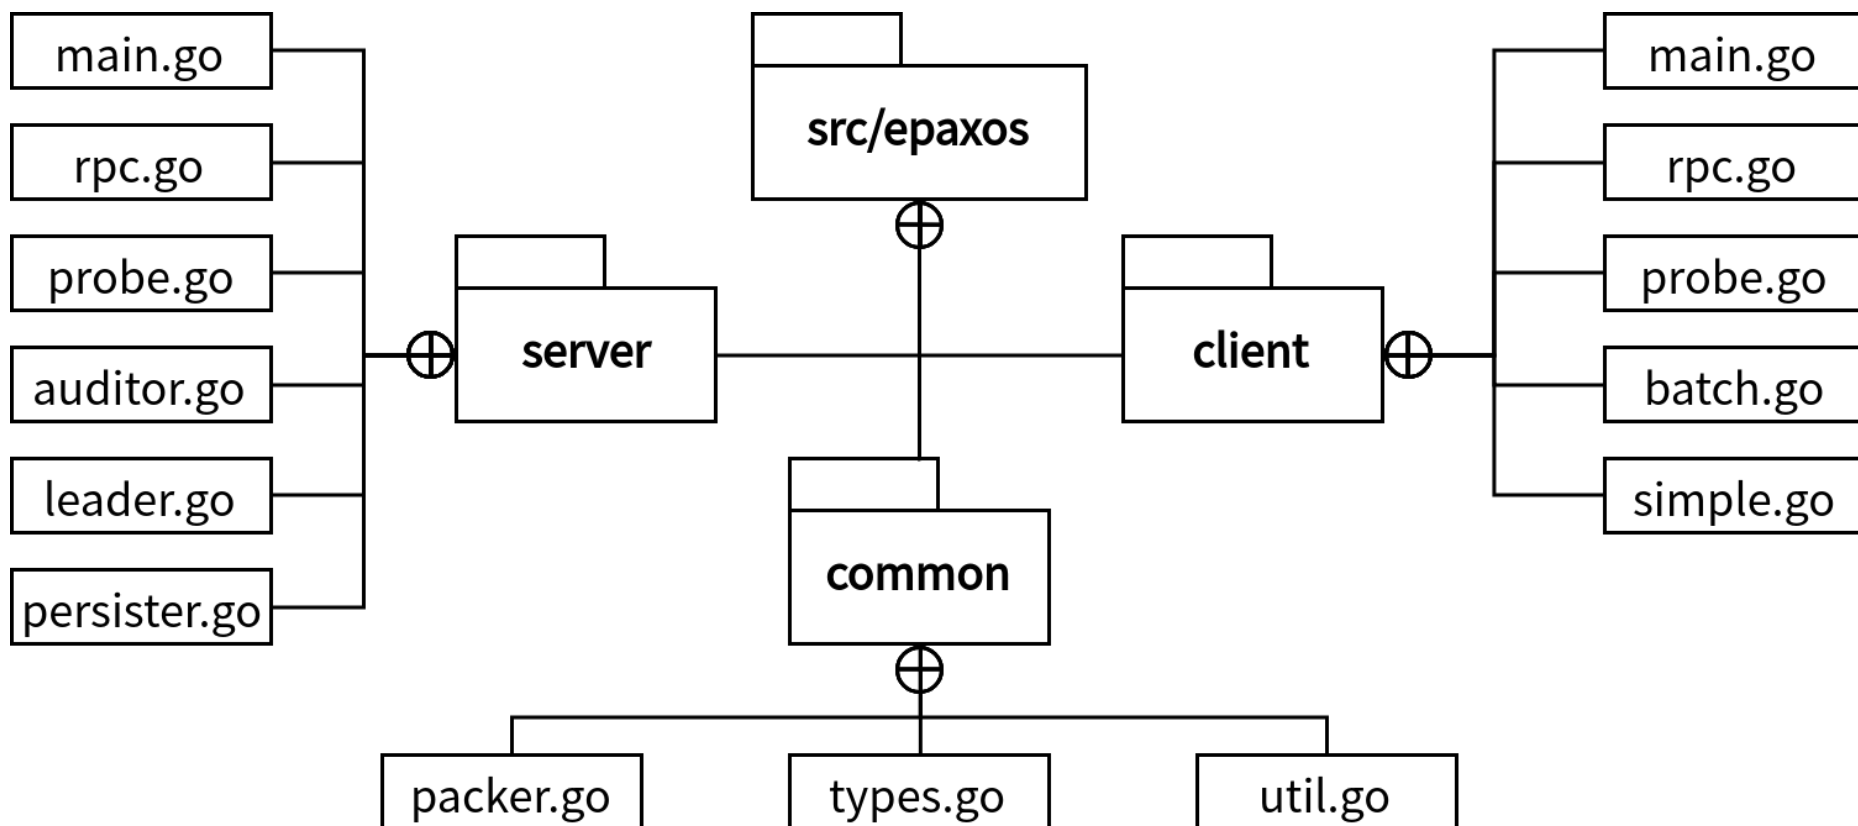
Logical View: Use Cases




Development View: Components



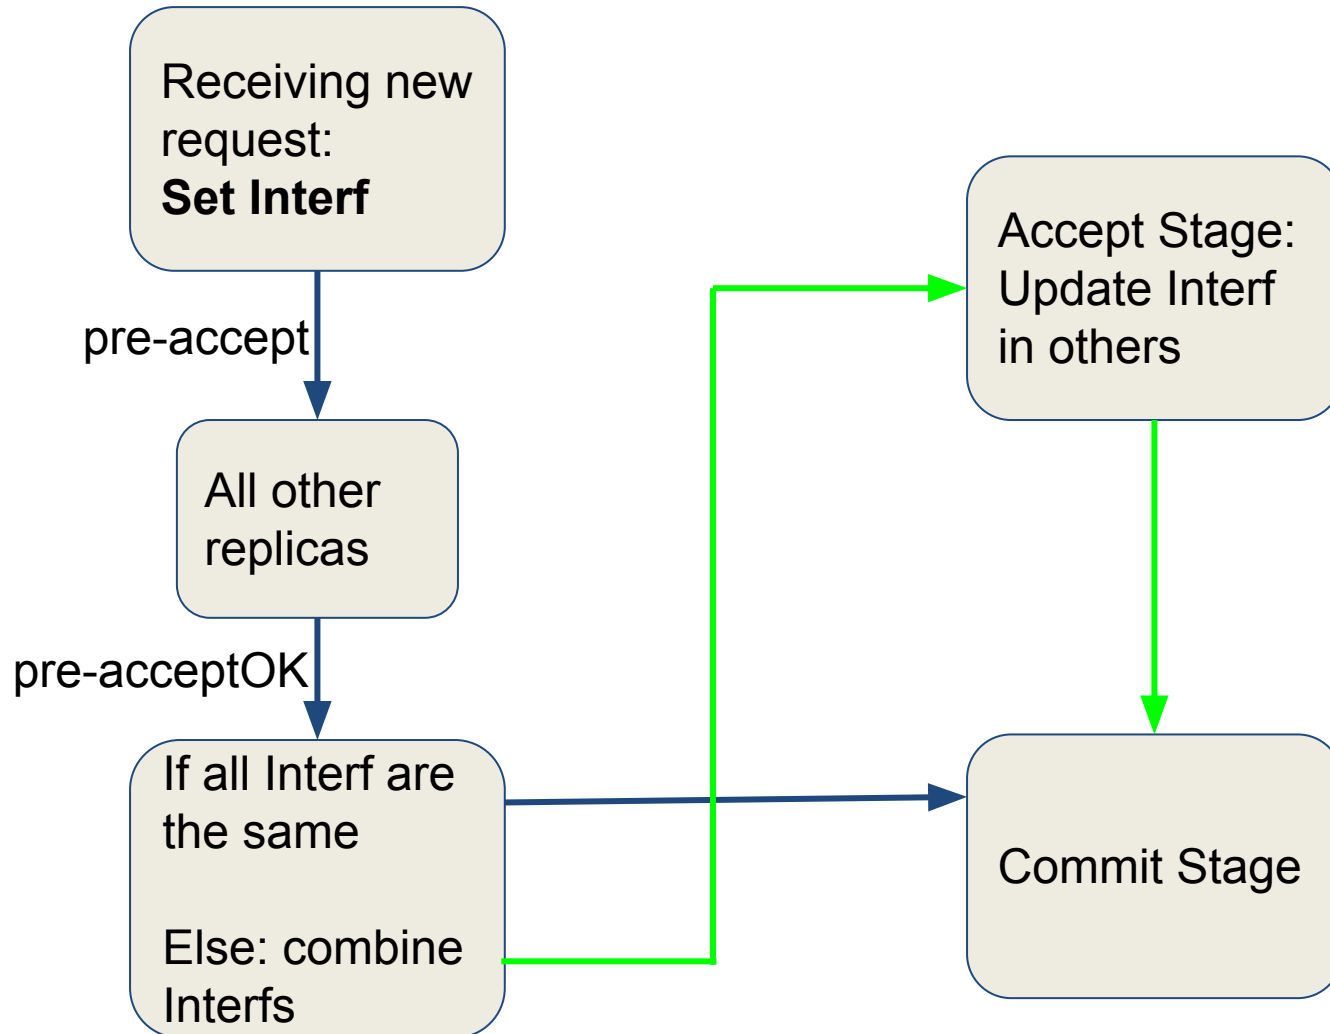
Development View: Packages



Physical View: Testing environment



Key Idea: How to make deps coherent



Key Idea: Deal with failure - Prepare

- One replica finds one potentially failed replica
 - Sends Prepare msg to all replicas
 - If replies include committed phase, run Commit phase
 - if replies include accepted phase, run Accept phase
 - if replies include $> N/2$ pre-accept phase, run Accept phase
 - Otherwise, run pre-accept phase