

# SILT

A Memory-Efficient, High-  
Performance Key-Value Store

# Introduction

# Motivation

---

- Problem: Read Amplification

# Motivation

---

- Problem: Read Amplification
- Solution(?): Keep index in DRAM

# Motivation

---

- Problem: Read Amplification
- Solution(?): Keep index in DRAM
- New Problem: Expense and lower capacity of DRAM

# Goals

---

# Goals

---

- Low # of flash reads per GET

# Goals

---

- Low # of flash reads per GET
- Sequential instead of Random Writes



# Goals

---

- Low # of flash reads per GET
- Sequential instead of Random Writes
- Memory efficient indexing

# Goals

---

- Low # of flash reads per GET
- Sequential instead of Random Writes
- Memory efficient indexing
- Computation efficient indexing

# Goals

---

- Low # of flash reads per GET
- Sequential instead of Random Writes
- Memory efficient indexing
- Computation efficient indexing
- Effective use of flash space

# Key Ideas

---

- Use three different KV stores
- Start in write-optimized store, move to memory-efficient store over time
- One flash read per lookup

# Implementation

# Overview


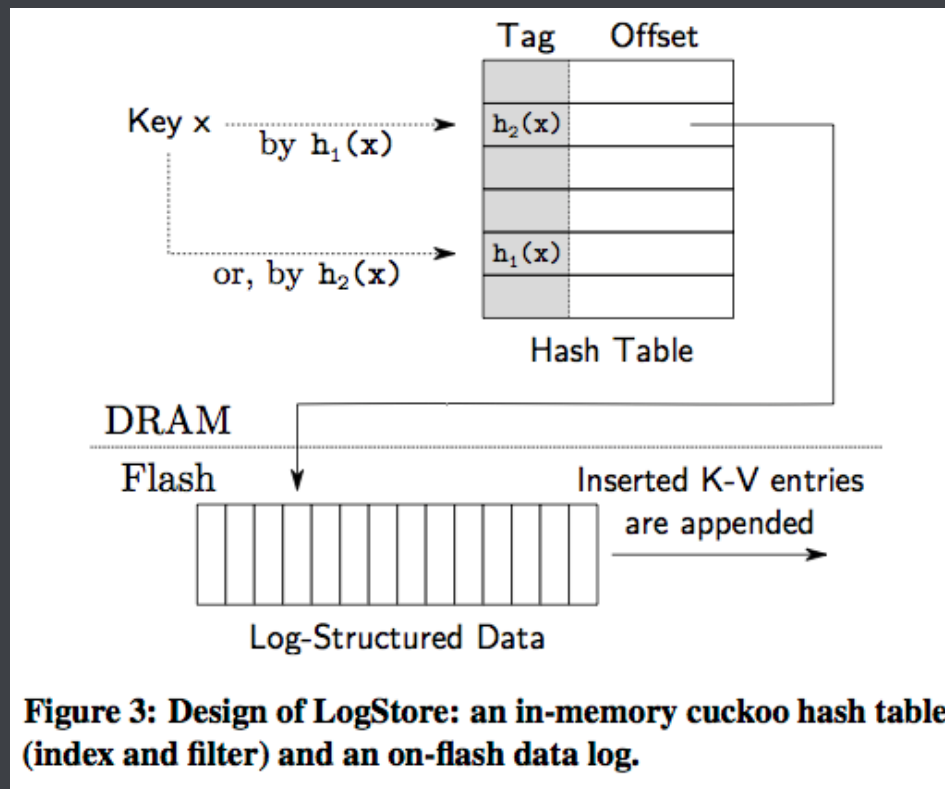


Figure 2: Architecture of SILT.

|               | SortedStore (§3.3)     | HashStore (§3.2) | LogStore (§3.1) |
|---------------|------------------------|------------------|-----------------|
| Mutability    | Read-only              | Read-only        | Writable        |
| Data ordering | Key order              | Hash order       | Insertion order |
| Multiplicity  | 1                      | $\geq 0$         | 1               |
| Typical size  | > 80% of total entries | < 20%            | < 1%            |
| DRAM usage    | 0.4 bytes/entry        | 2.2 bytes/entry  | 6.5 bytes/entry |

Table 2: Summary of basic key-value stores in SILT.

# Log Store



- Cuckoo hash table index + flash log
- Ops written to log sequentially
- Convert to HashStore when index is full

# Hash Store

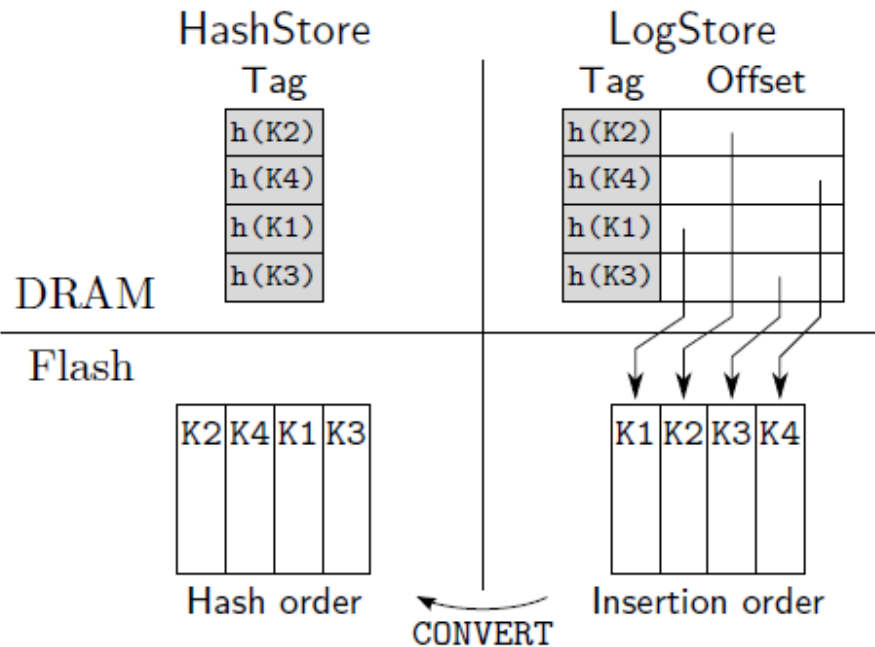



Figure 4: Convert a LogStore to a HashStore. Four keys  $K1$ ,  $K2$ ,  $K3$ , and  $K4$  are inserted to the LogStore, so the layout of the log file is the insert order; the in-memory index keeps the offset of each key on flash. In HashStore, the on-flash data forms a hash table where keys are in the same order as the in-memory filter.

- DRAM tags preserve order on flash
- Lower memory overhead than LS
- Merge multiple into SortedStore at once



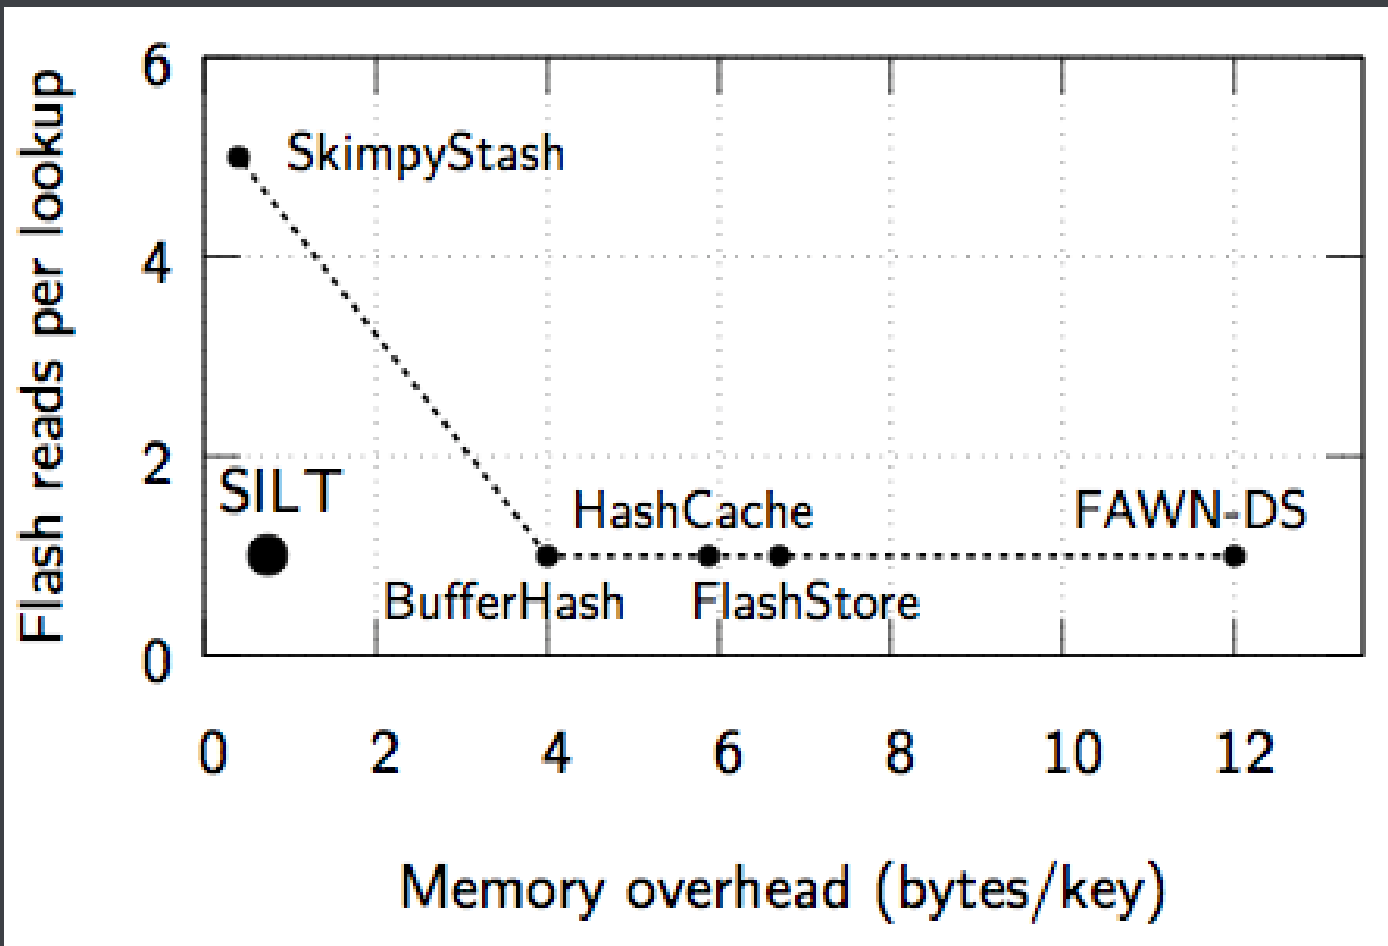
# Sorted Store

- DRAM trie index + sorted KVs on flash
- Index compressed with entropy coding
- Stores ~80% of data




# Analysis

# Evaluation



SILT vs. other KV stores

# Tradeoffs



Improving one metric will cost you in another

# Takeaways

---

- Composition of existing systems can make a better system
- Avoiding amplification is hard
- System designers need to choose which metric to focus on

Thank You!