


RIPQ

Advanced Photo Caching on Flash for Facebook

Linpeng Tang, Qi Huang, Wyatt Lloyd, Sanjeev Kumar, Kai Li

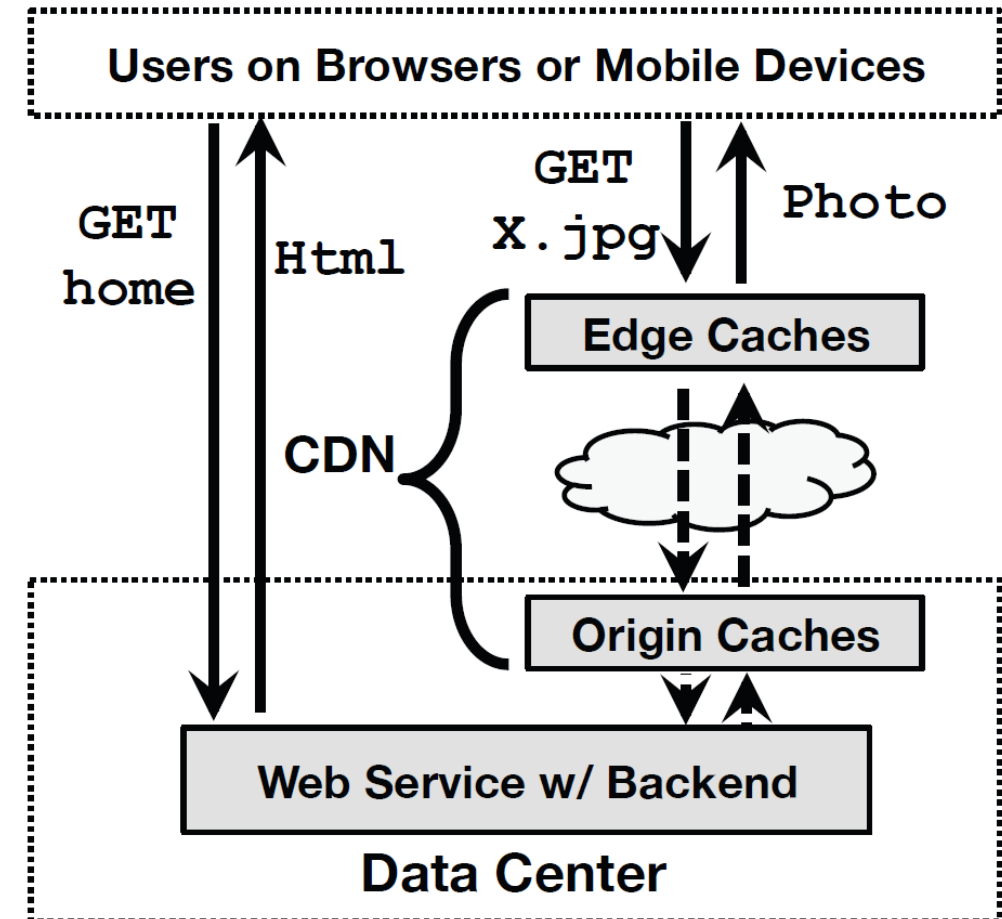
Motivation

- Need to improve **hit ratios** in Facebook's **flash-based** distributed photo caches
- Advanced caching algorithms generate many **small random writes**
 - Frequent garbage collection in FTL \Rightarrow *high write amplification, shorter device lifespan*
 - Small writes \Rightarrow *low throughput*
- Result: Facebook used FIFO caches
 - Frequent cache misses \Rightarrow *many back-end I/Os*



Background: Cache Layers

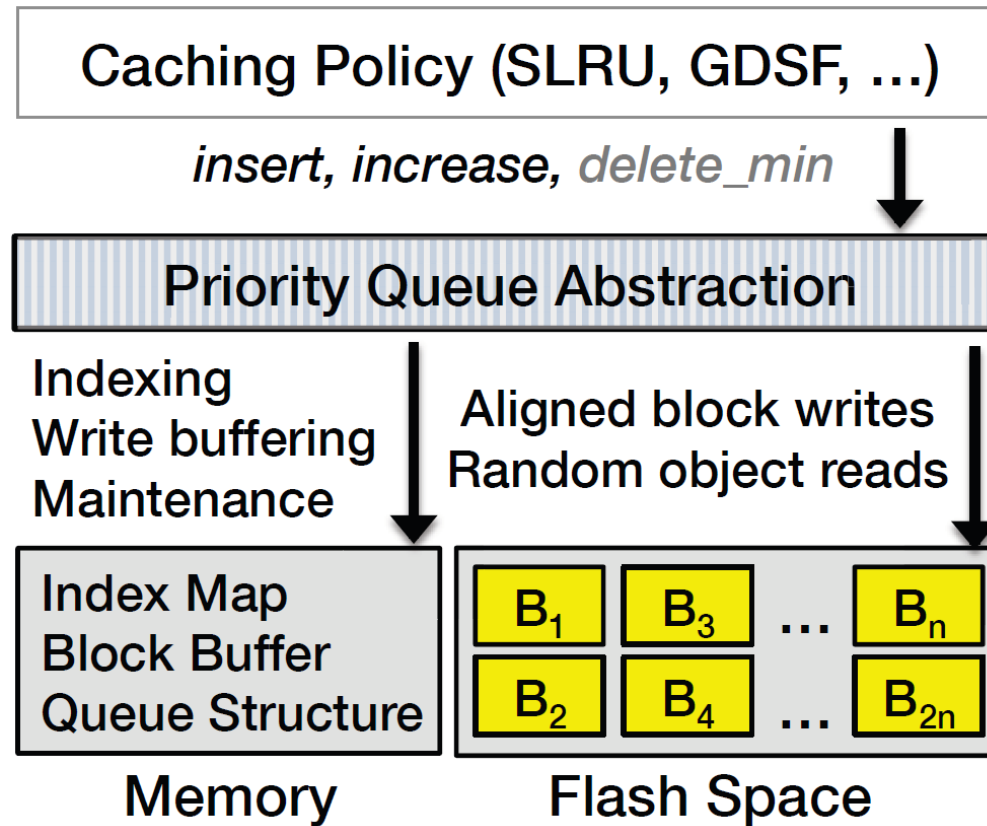
- **Edge caches** close to users
 - Reducing *traffic to datacenters*: caches should maximize **byte-wise** hit ratio
- **Origin cache** in datacenters
 - Reducing *back-end I/Os*: caches should maximize **object-wise** hit ratio
- Design a system to support both?



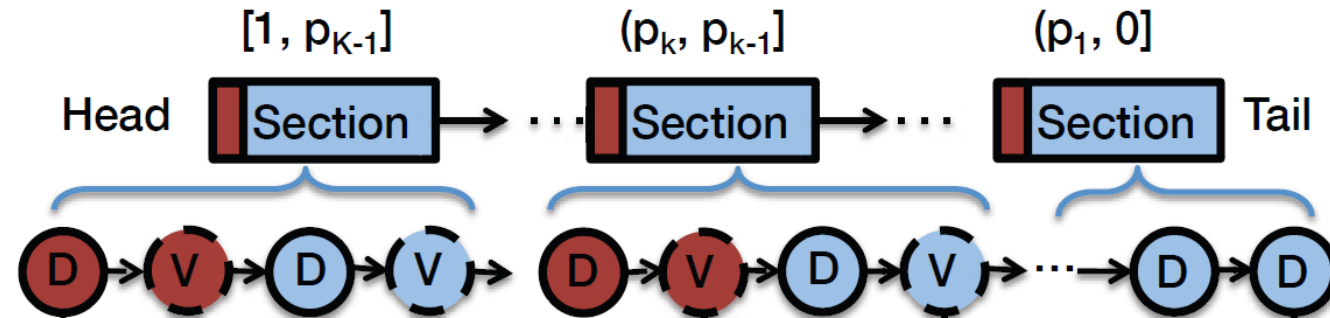
Restricted Insertion Priority Queue (RIPQ)

- General framework to support *many* advanced caching techniques
 - *Segmented-LRU (SLRU)*: better *byte-wise* hit ratio (edge caches)
 - *Greedy-Dual-Size-Frequency (GDSF)*: better *object-wise* hit ratio (origin cache)
- **Priority queue** abstraction – but only an **approximation!**
 - $\text{insert}(x, p)$: insert object x with priority p
 - $\text{increase}(x, p)$: increase priority of object x to p
 - $\text{delete-min}(x, p)$: delete object with lowest priority

System Architecture




Queue Sections




- Divide *priority space* $[0,1]$ into K **sections**, ex. $[1,0.7]$, $(0.7,0.3]$, $(0.3,0]$
 - Sections composed of **blocks** (flash data unit)
- **Insert** new objects at the *head* of a section (based on priority)
 - All objects in *same or lower section* are **implicitly demoted** in the queue
 - Restricted to K insertion points: *accuracy vs. space* tradeoff

Device Blocks



- Each section has *one active device block* and *many sealed device blocks*
 - *Active block* accepts and **buffers** object insertions in **memory**
 - When active block is full: **flush** contents to **flash**, turn into a *sealed block*
- **Index map** associates objects' keys with location (flash/RAM, block ID, ...)

Virtual Blocks

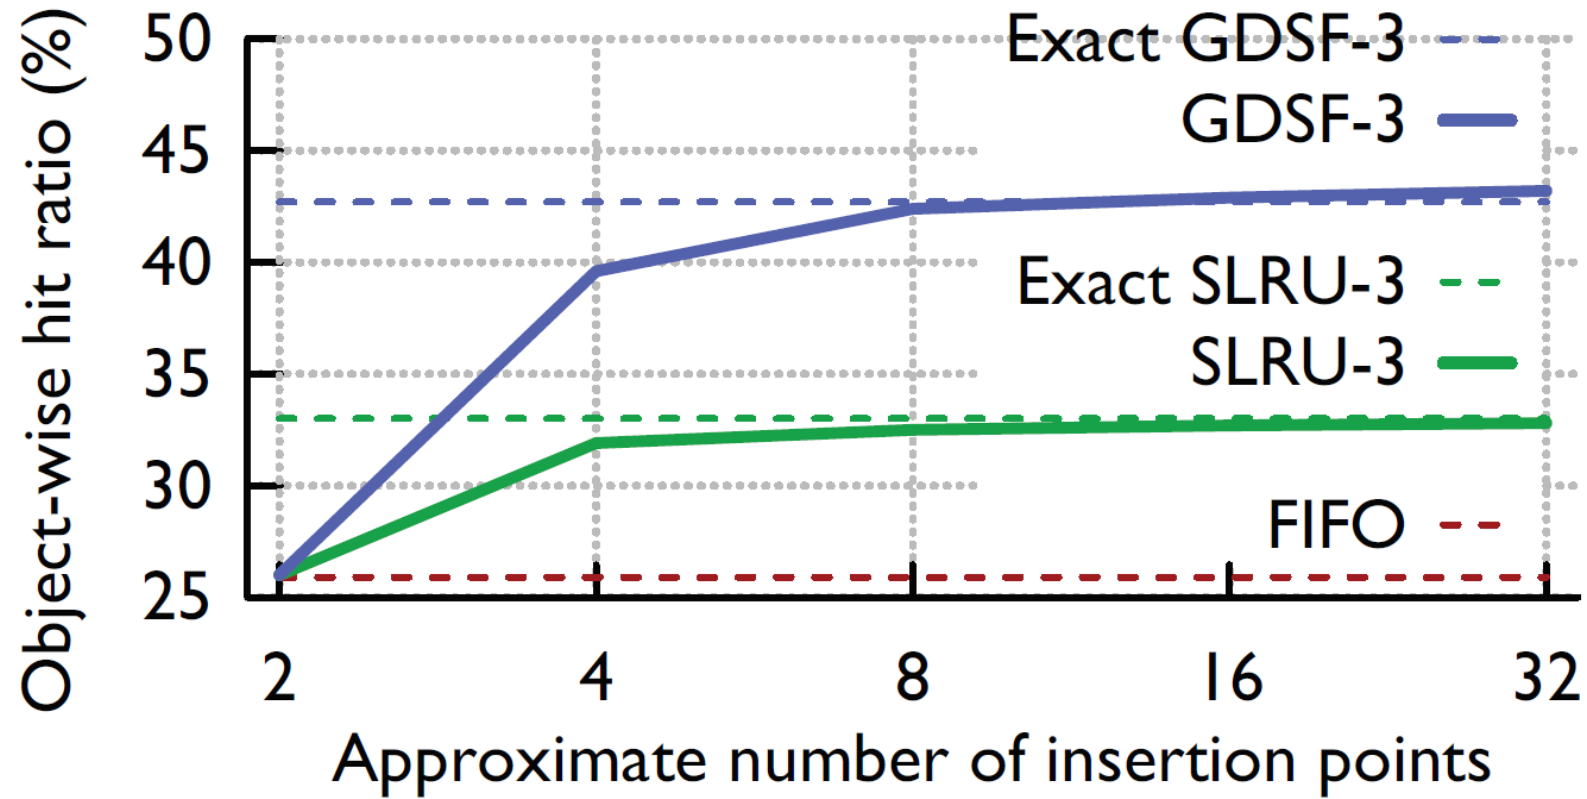


- To handle **priority increases**, use **lazy updates** instead of *duplicating* an object
- Each section also has *one active virtual block* and *many sealed virtual blocks*
 - Updates only add a *pointer* in the virtual block, and set a *virtual block ID* in the *index map*
 - Upon *eviction* of a *device block*, objects with a *virtual block ID* are *re-inserted* into the corresponding *active device block*

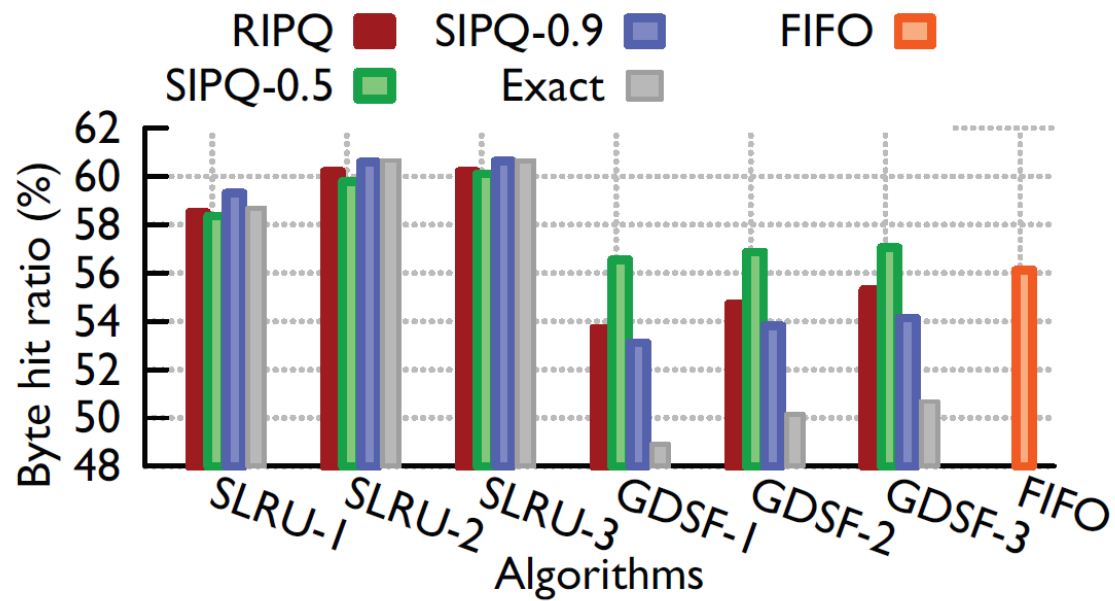
Single Insertion Priority Queue (SIPQ)

- Simplifies RIPQ for use in memory-constrained settings
- Good results for simpler algorithms (ex. LRU)

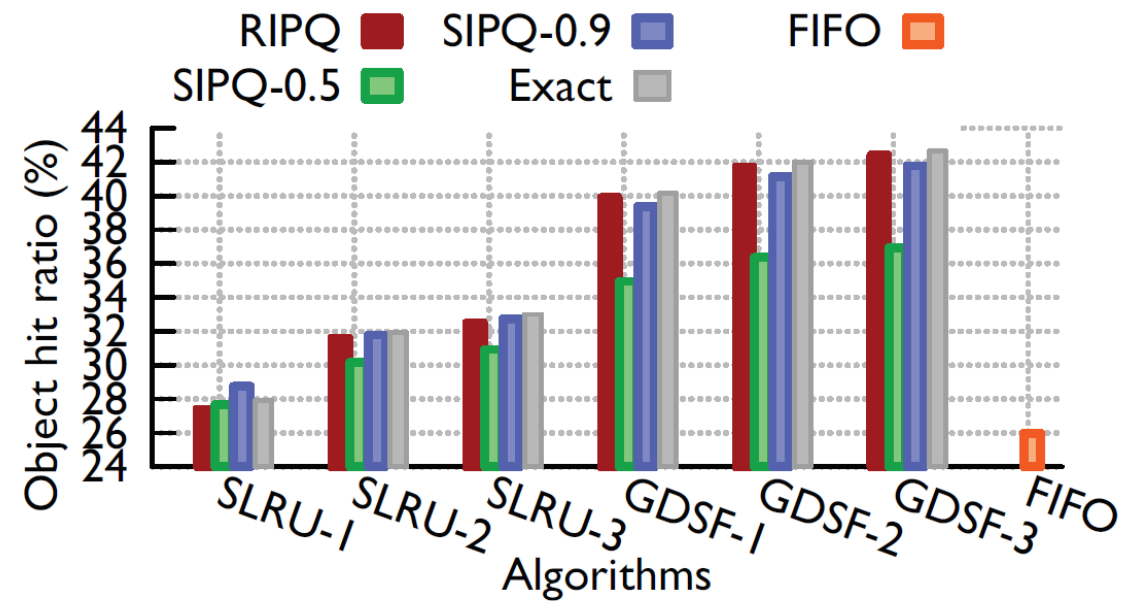
Evaluation: Fidelity



Evaluation: Hit Ratios

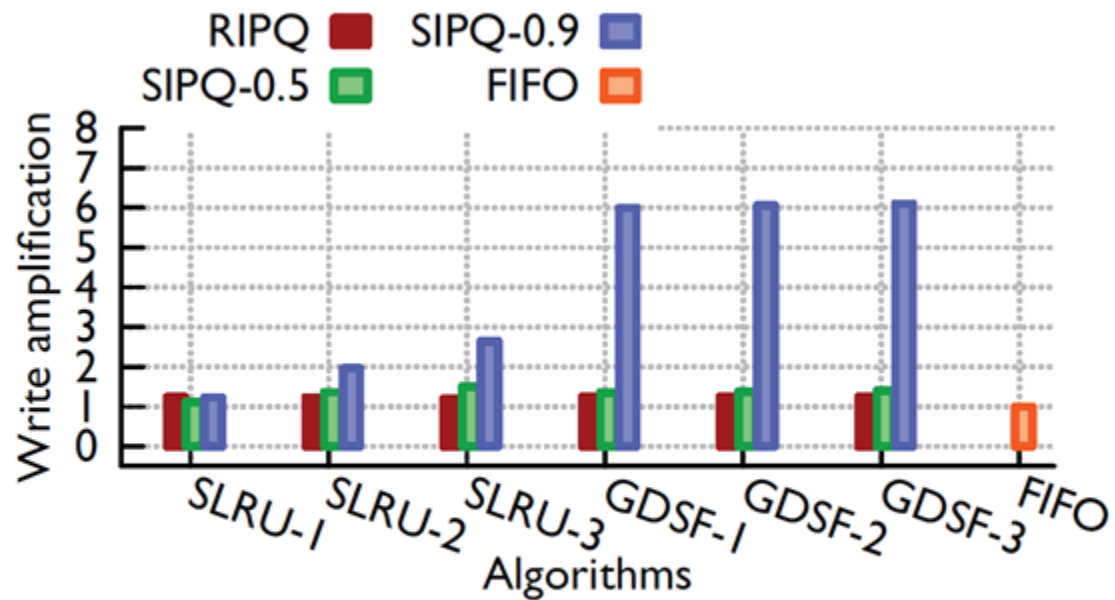


Byte-wise hit ratio (Edge)

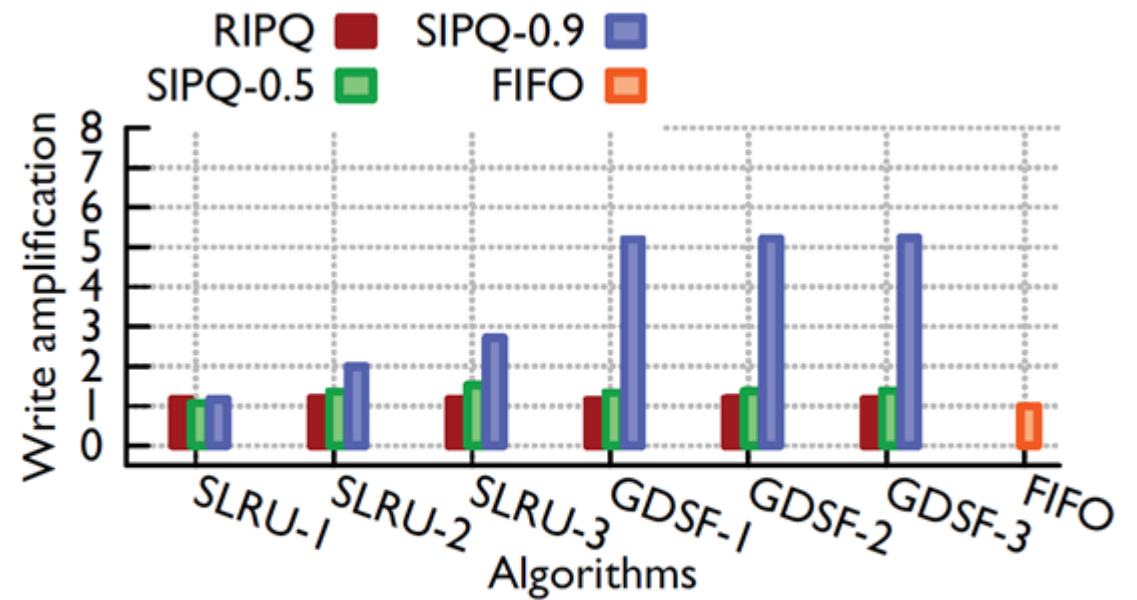


Object-wise hit ratio (Origin)

Evaluation: Write Amplification



Write amplification (Edge)



Write amplification (Origin)